

Exercise Set 1 | PML Fall 2022

Unless indicated otherwise the exercises refer to "Pattern Recognition and Machine Learning" by CM Bishop (Springer) 2006.

Theory

Basic Probability Theory

1.10)

Let x, y be independent random variables, then the linear property of E is: the expectation is

$$\begin{aligned} E[x + z] &= \sum_x \sum_z p(x, z)(x + z) = \sum_x \sum_z p(x)p(z)(x) + \sum_x \sum_z p(x)p(z)(z) \\ &= \sum_x p(x)x \sum_z p(z) + \sum_z p(z)z \sum_x p(x) = E[x] \times 1 + E[z] \times 1, \end{aligned}$$

$$\begin{aligned} \text{and variance as } Var[x + z] &= E[(x + z)^2] - E[(x + z)]^2 \\ &= E[x^2 + 2xz + z^2] - (E[x] + E[z])^2 \\ &= E[x^2] + 2E[x]E[z] + E[z^2] - E[x]^2 - 2E[x]E[z] - E[z]^2 \\ &= E[x^2] - E[x]^2 + E[z^2] - E[z]^2 = Var[x] + Var[z]. \end{aligned}$$

2.8)

Consider variables x, y jointly distributed with $p(x, y)$, then:

$$E_y[E_x[x|y]] = E_y[\sum_x p(x|y)x] = \sum_y p(y) \sum_x p(x|y)x = \sum_x p(x)x = E[x].$$

For the variance it is the case that:

$$\begin{aligned} E_y[V_x[x|y]] &= E_y[E_x[x^2|y] - E_x[x|y]^2] = E_y[E_x[x^2|y]] - E_y[E_x[x|y]^2] = E_x[x^2] - E_y[E_x[x|y]^2] \\ \text{and } V_x[E_x[x|y]] &= E_y[E_x[x|y]^2] - E_y[E_x[x|y]]^2 \text{ therefore it follows that} \\ V[x] &= E_x[x^2] - E_y[E_x[x|y]^2] + E_y[E_x[x|y]^2] - E_y[E_x[x|y]]^2 = E_x[x^2] - E_x[x]^2. \end{aligned}$$

2.20)

A p.d. real valued matrix Σ defined in the quadratic form as $\mathbf{a}^T \Sigma \mathbf{a}$ is positive for any real value of vector \mathbf{a} , because (given symmetry and decomposition) and following matrix notation for (2.45) and (2.48) s.t. $U^T \mathbf{a} = \mathbf{y} \neq 0$ then: $\mathbf{a}^T \Sigma \mathbf{a} = \mathbf{a}^T (U^T \Lambda U) \mathbf{a} = \mathbf{y}^T \Lambda \mathbf{y} = \sum_i \lambda_i y_i^2 > 0$. Here

matrix symmetry and real values allow us to use *spectral decomposition*.

Otherwise using (2.45) : $\mathbf{a}^T \Sigma \mathbf{a} = \sum_i \lambda_i \mathbf{u}_i^T \Sigma \sum_j \lambda_j \mathbf{u}_j = \sum_{ij} \lambda_i \lambda_j \mathbf{u}_i^T \Sigma \mathbf{u}_j$ and as orthonormal basis we can write $\Sigma \mathbf{u}_j = \phi_j \mathbf{u}_j$ for eigenvalues ϕ , then $\Rightarrow \sum_{ij} \lambda_i \lambda_j \phi_j \mathbf{u}_i^T \mathbf{u}_j = \sum_i \lambda_i^2 \phi_i > 0$.

See Matrix Cookbook (293) section 5.3.1 applicable to real symmetric matrices.

8.3)

Evaluating the joint distribution over three binary variables

true

```
begin
    using DataFrames
    table = DataFrame(
        a=[0,0,0,0,1,1,1,1],
        b=[0,0,1,1,0,0,1,1],
        c=[0,1,0,1,0,1,0,1],
        p_abc=[0.192,0.144,0.048,0.216,0.192,0.064,0.048,0.096])

    @show table
    println()

    function test_conditional(a, b, c, table::DataFrame=table)
        p_c = sum(table[table.c .== c, :p_abc])
        p_a_c = sum(filter(row -> row.a == a && row.c == c, table).p_abc)/p_c
        p_b_c = sum(filter(row -> row.b == b && row.c == c, table).p_abc)/p_c
        p_a_b_c = sum(filter(row -> row.a == a && row.b==b && row.c==c,
            table).p_abc)/p_c
        println("p(a=$a,b=$b|c=$c)=$p_a_b_c =
            $p_a_c*$p_b_c=p(a=$a|c=$c)p(b=$b|c=$c)")
        #@assert isequal(p_a_b_c, p_a_c * p_b_c)
        return isequal(p_a_b_c, p_a_c * p_b_c)
    end

    p_a = sum(table[table.a .== 1, :p_abc])
    p_b = sum(table[table.b .== 1, :p_abc])
    # a and b are marginally dependent:
    @assert p_a * p_b != p_a + p_b
    # test all for conditional independence
    map(test_conditional, table.a, table.b, table.c) |> all
end
```

```
table = 8x4 DataFrame
```

Row	a Int64	b Int64	c Int64	p_abc Float64
1	0	0	0	0.192
2	0	0	1	0.144
3	0	1	0	0.048
4	0	1	1	0.216
5	1	0	0	0.192
6	1	0	1	0.064
7	1	1	0	0.048
8	1	1	1	0.096

```
p(a=0,b=0|c=0)=0.4 = 0.5*0.8=p(a=0|c=0)p(b=0|c=0)
p(a=0,b=0|c=1)=0.2769230769230769 = 0.6923076923076923*0.39999999999999997
=p(a=0|c=1)p(b=0|c=1)
p(a=0,b=1|c=0)=0.1 = 0.5*0.2=p(a=0|c=0)p(b=1|c=0)
p(a=0,b=1|c=1)=0.41538461538461535 = 0.6923076923076923*0.6=p(a=0|c=1)p(b=
1|c=1)
p(a=1,b=0|c=0)=0.4 = 0.5*0.8=p(a=1|c=0)p(b=0|c=0)
p(a=1,b=0|c=1)=0.12307692307692307 = 0.3076923076923077*0.39999999999999999
7=p(a=1|c=1)p(b=0|c=1)
p(a=1,b=1|c=0)=0.1 = 0.5*0.2=p(a=1|c=0)p(b=1|c=0)
p(a=1,b=1|c=1)=0.18461538461538463 = 0.3076923076923077*0.6=p(a=1|c=1)p(b=
1|c=1)
```

8.4)

Show $p(a, b, c) = p(a)p(c|a)p(b|c)$.

Remark: $p(a, b, c) = p(a) \frac{p(c,a)}{p(a)} \frac{p(b,c)}{p(c)} = p(c, a) \frac{p(b,c)}{p(c)}$.

true

```

begin
  function test_joint(a,b,c,table::DataFrame=table)
    p_c = sum(table[table.c .== c, :p_abc])
    p_c_a = sum(filter(row -> row.c == c && row.a == a, table).p_abc)
    p_b_c = sum(filter(row -> row.b == b && row.c == c, table).p_abc)/p_c
    p_a_b_c = sum(filter(row -> row.a == a && row.b==b && row.c==c,
      table).p_abc)
    joint = p_c_a*p_b_c
    println("p(a=$a,b=$b,c=$c)=$p_a_b_c = $joint=p(c,a)*p(b|c)/p(c)")
    # @show p_c
    # @show p_c_a
    # @show p_b_c
    # @show p_a_b_c
    return isequal(p_c_a*p_b_c, p_a_b_c)
  end
  map(test_joint, table.a, table.b, table.c) |> all
end

```

```

p(a=0,b=0,c=0)=0.192 = 0.192=p(c,a)*p(b|c)/p(c)
p(a=0,b=0,c=1)=0.144 = 0.144=p(c,a)*p(b|c)/p(c)
p(a=0,b=1,c=0)=0.048 = 0.048=p(c,a)*p(b|c)/p(c)
p(a=0,b=1,c=1)=0.216 = 0.216=p(c,a)*p(b|c)/p(c)
p(a=1,b=0,c=0)=0.192 = 0.192=p(c,a)*p(b|c)/p(c)
p(a=1,b=0,c=1)=0.064 = 0.064=p(c,a)*p(b|c)/p(c)
p(a=1,b=1,c=0)=0.048 = 0.048=p(c,a)*p(b|c)/p(c)
p(a=1,b=1,c=1)=0.096 = 0.096=p(c,a)*p(b|c)/p(c)

```

8.9)

D-separation, we recall:

- i) $[A] \rightarrow [X] \rightarrow [B]$; $[A] \leftarrow [X] \rightarrow [B]$ then $A \perp\!\!\!\perp B | X$, if X is in the conditioning set.
- ii) $[A] \rightarrow [X] \leftarrow [B]$ $A \perp\!\!\!\perp B$ if X not in conditioning set.

For the Markov Blanket see the definition from Bishop (Fig.8.26) and Lecture notes. The answer follows directly from the definition in these materials.

8.11)

$[B] \rightarrow [G] \leftarrow [F]$

\downarrow
 $[D]$

0.10962566844919786

```

begin
    p_b1=0.9
    p_b0=1-p_b1
    p_f1=0.9
    p_f0=0.1
    p_g0=0.315
    p_g0_f0=0.81
    p_f0_g0=0.257
    p_g1_b1_f1=0.8
    p_g0_b1_f1=1-p_g1_b1_f1
    p_g1_b1_f0=0.2
    p_g0_b1_f0=1-p_g1_b1_f0
    p_g1_b0_f1=0.2
    p_g0_b0_f1=1-p_g1_b0_f1
    p_g1_b0_f0=0.1
    p_g0_b0_f0=1-p_g1_b0_f0
    p_d1_g1 = 0.9
    p_d0_g1 = 1-p_d1_g1
    p_d0_g0 = 0.9
    p_d1_g0 = 1-p_d0_g0
    ## i)
    # step 1: compute likelihood
    p_d0_f0 = (p_d0_g0*p_g0_b0_f0*p_b0+p_d0_g1*p_g1_b0_f0*p_b0) +
    (p_d0_g0*p_g0_b1_f0*p_b1+p_d0_g1*p_g1_b1_f0*p_b1)
    # step 2: compute marginal
    p_d0 =
    (p_d0_g0*p_g0_b0_f0*p_b0*p_f0+p_d0_g1*p_g1_b0_f0*p_b0*p_f0+p_d0_g0*p_g0_b1_f
    0*p_b1*p_f0+p_d0_g1*p_g1_b1_f0*p_b1*p_f0)+
    (p_d0_g0*p_g0_b0_f1*p_b0*p_f1+p_d0_g1*p_g1_b0_f1*p_b0*p_f1+p_d0_g0*p_g0_b1_f
    1*p_b1*p_f1+p_d0_g1*p_g1_b1_f1*p_b1*p_f1) # (F=0) + (F=1)
    # step 3: Bayes Rule
    @show p_f0_d0 = (p_d0_f0*p_f0)/p_d0
    ## ii)
    # step 1: compute likelihood
    p_d0_b0_f0 = (p_d0_g0*p_g0_b0_f0*p_b0) + (p_d0_g1*p_g1_b0_f0*p_b0)
    # step 2: compute marginal
    p_d0_b0 = p_d0_g0*p_g0_b0_f0*p_b0*p_f0 + p_d0_g1*p_g1_b0_f0*p_b0*p_f0 +
    p_d0_g0*p_g0_b0_f1*p_b0*p_f1 + p_d0_g1*p_g1_b0_f1*p_b0*p_f1
    # step 3: Bayes Rule
    @show p_f0_d0_b0 = (p_d0_b0_f0*p_f0)/p_d0_b0
end

```

```

p_f0_d0 = (p_d0_f0 * p_f0) / p_d0 = 0.212500000000000008
p_f0_d0_b0 = (p_d0_b0_f0 * p_f0) / p_d0_b0 = 0.10962566844919786

```

Programming Exercise

TaskLocalRNG()

```
• begin
•     using Random
•     using Distributions
•     using LinearAlgebra
•     using PlutoUI
•     Random.seed!(42)
• end
```

Solve for D: ☐

```
(2×2 Matrix{Float64}[:, 0.1)
 1.0  0.0
 0.0  1.0
```

```
• begin
•     N=2
•     Σ_x = convert(Array{Float64}, collect(I(N)))
•     σ = 0.1 # Option: A,B,C
•     if option
•         Σ_x[1,1] = 0.1 # Option for D)
•         # σ = 0.01 # Even more options here...
•     end
•     Σ_x, σ
• end
```

A)

Create dataset for regression.:

```
(20×2 Matrix{Float64}: , [0.3581, 0.705565, -0.244711, 1.2944, 0.472981, 2.10565, 0.
-0.954618 -0.61694
 0.888373  1.70221
-0.307589 -0.595113
 0.240871  1.34142
-0.312024  0.220427
 0.199745  2.21349
-1.41501   -0.83954
  ⋮
-0.120277 -0.289603
-1.14098   0.220411
 0.520437  0.765463
 1.37368   1.73886
 0.576634 -0.0593636
-1.16674   -0.973451
```

```
• begin
•     sample_size=20
•     dim = 2
•     μ = [0., 0.]
•     N_x = MvNormal(μ, Σ_x)
•     X = reshape(rand(N_x, sample_size), (sample_size, dim))
•     θ = [-1,1]
•     y = (X*θ) .+ rand(Normal(0, σ), sample_size)
•     X, y
• end
```

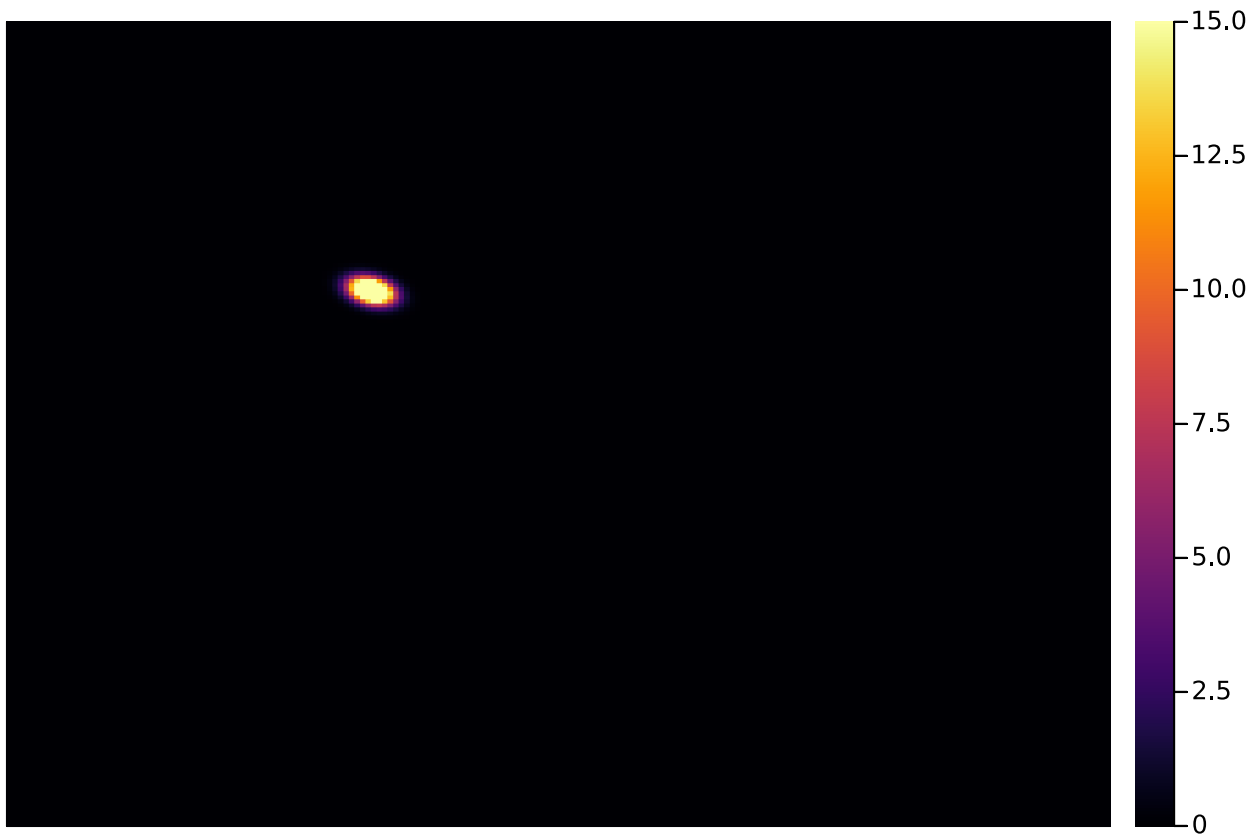
B)

Compute posterior mean and variance - see script Eq. (55),(56):

```
2×2 Matrix{Float64}:
 0.00583072 -0.00132868
-0.00132868  0.00395157
```

```
• begin
•     μ_θ = X'*inv(σ.*I + X*X') * y # Eq.(55)
•     Σ_θ = I - X'*inv(σ.*I + X*X') * X # Eq.(56)
•     @show μ_θ
•     @show Σ_θ
• end
```

```
μ_θ = [-1.0226511074024365, 0.9979297993773675]
Σ_θ = [0.005830718405949509 -0.001328680661916788; -0.0013286806619201591
0.003951568963384555]
```



```

• begin
•   # plot PDF
•   using Base.Iterators
•   using Plots
•   θ_range = range(-3, stop=3, length=200)
•   xx = collect(product(θ_range, θ_range))
•   XX = [x[i] for x in [xx...], i ∈ 1:2]
•   N_θ = MvNormal(μ_θ, Hermitian(Σ_θ)) # force it to assume p.d. cov.
•   values = reshape(pdf(N_θ, XX'), (size(θ_range)[1], size(θ_range)[1]))
•   heatmap(values', clim=(0, 15), xaxis=nothing, yaxis=nothing)
• end

```

C)

Compute Variance of posterior predictive



```
• begin
•     variance = sum((XX*Σ_θ).*XX, dims=2) .+ σ
•     variance = reshape(variance, (size(θ_range)[1], size(θ_range)[1]))
•     heatmap(variance', clim=(0., 0.4), xaxis=nothing, yaxis=nothing)
• end
```